

```
[3]: import pandas as pd
```

```
[4]: ## Data Ingestion step
df=pd.read_csv('https://raw.githubusercontent.com/krishnaik86/FS05Regression/main/notebooks/data/gemstone.csv')
df.head()
```

```
[4]:
```

	id	carat	cut	color	clarity	depth	table	x	y	z	price
0	0	1.52	Premium	F	VS2	62.2	58.0	7.27	7.33	4.55	13619
1	1	2.03	Very Good	J	S12	62.0	58.0	8.06	8.12	5.05	13387
2	2	0.70	Ideal	G	VS1	61.2	57.0	5.69	5.73	3.90	2772
3	3	0.32	Ideal	G	VS1	61.6	56.0	4.38	4.41	2.71	666
4	4	1.70	Premium	G	VS2	62.6	59.0	7.65	7.61	4.77	14453

```
[5]: df.isnull().sum()
```

```
[5]:
```

id	0
carat	0
cut	0
color	0
clarity	0
depth	0
table	0
x	0
y	0
z	0
price	0

dtype: int64

```
[6]: ### No missing values present in the data
```

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 193573 entries, 0 to 193572
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   id           193573 non-null  int64
1   carat        193573 non-null  float64
2   cut          193573 non-null  object
3   color        193573 non-null  object
4   clarity      193573 non-null  object
5   depth        193573 non-null  float64
6   table        193573 non-null  float64
7   x            193573 non-null  float64
8   y            193573 non-null  float64
9   z            193573 non-null  float64
10  price        193573 non-null  int64
dtypes: float64(6), int64(2), object(3)
memory usage: 16.2+ MB
```

```
[7]: df.head()
```

```
[7]:
```

	id	carat	cut	color	clarity	depth	table	x	y	z	price
0	0	1.52	Premium	F	VS2	62.2	58.0	7.27	7.33	4.55	13619
1	1	2.03	Very Good	J	SI2	62.0	58.0	8.06	8.12	5.05	13387
2	2	0.70	Ideal	G	VS1	61.2	57.0	5.69	5.73	3.50	2772
3	3	0.32	Ideal	G	VS1	61.6	56.0	4.38	4.41	2.71	666
4	4	1.70	Premium	G	VS2	62.6	59.0	7.65	7.61	4.77	14453

```
[8]: ## Lets drop the id column
df=df.drop(labels=['id'],axis=1)
df.head()
```

```
[8]:
```

	carat	cut	color	clarity	depth	table	x	y	z	price
0	1.52	Premium	F	VS2	62.2	58.0	7.27	7.33	4.55	13619
1	2.03	Very Good	J	SI2	62.0	58.0	8.06	8.12	5.05	13387
2	0.70	Ideal	G	VS1	61.2	57.0	5.69	5.73	3.50	2772
3	0.32	Ideal	G	VS1	61.6	56.0	4.38	4.41	2.71	666
4	1.70	Premium	G	VS2	62.6	59.0	7.65	7.61	4.77	14453

```
[9]: ## check for duplicated records
df.duplicated().sum()
```

```
[9]: 0
```

```
[10]: ## segregate numerical and categorical columns
numerical_columns=df.columns[df.dtypes!='object']
categorical_columns=df.columns[df.dtypes=='object']
print("Numerical columns:",numerical_columns)
print("Categorical columns:",categorical_columns)
```

```
Numerical columns: Index(['carat', 'depth', 'table', 'x', 'y', 'z', 'price'], dtype='object')
Categorical columns: Index(['cut', 'color', 'clarity'], dtype='object')
```

```
[11]: df[categorical_columns].describe()
```

```
[11]:
```

	cut	color	clarity
count	193573	193573	193573
unique	5	7	8
top	Ideal	G	SI1
freq	92454	44391	53272

```
[2]: df['cut'].value_counts()
```

```
[2]: cut
Ideal      91454
Premium    49018
Very Good  37566
Good       11621
Fair       2821
Name: count, dtype: int64
```

```
[3]: df['color'].value_counts()
```

```
[3]: color
G      44391
E      35869
F      34258
H      38799
D      24286
I      17514
J       6456
Name: count, dtype: int64
```



```
[4]: df['clarity'].value_counts()
```



```
[4]: clarity
SI1      53272
VS2      48827
VS1      38669
SI2      38484
VVVS2    15762
VVVS1    18628
IF        4219
I1         512
Name: count, dtype: int64
```

```
import seaborn as sns
```



```
11
```

```
import matplotlib.pyplot as plt
plt.figure(figsize=(8,6))
x=0
for i in numerical_columns:
    sns.histplot(data=df,x=i,kde=True)
    print("\n")
plt.show()
```







table




y



price

```
In [27]: ## correlation
sns.heatmap(df.corr(),annot=True)
```

Out[27]: <Axes: >



```
In [28]: df.head()
```

```
Out[28]:
```

	carat	cut	color	clarity	depth	table	x	y	z	price
0	1.52	Premium	F	VS2	62.2	58.0	7.27	7.33	4.55	13619
1	2.03	Very Good	J	S2	62.0	58.0	8.06	8.12	5.05	13387
2	0.70	Ideal	G	VS1	61.2	57.0	5.69	5.73	3.50	2772
3	0.32	Ideal	G	VS1	61.5	56.0	4.38	4.41	2.71	666
4	1.70	Premium	G	VS2	62.5	59.0	7.65	7.61	4.77	14453


```
10) df['cut'].unique()
11) array(['Premium', 'Very Good', 'Ideal', 'Good', 'Fair'], dtype=object)
```

```
12) cut_map={'Fair':1,'Good':2,'Very Good':3,'Premium':4,'Ideal':5}
```

```
13) df['clarity'].unique()
```

```
14) array(['VS2', 'SI2', 'VS1', 'SI1', 'IF', 'VVS2', 'WS1', 'I1'],
      dtype=object)
```

```
15) clarity_map = {"I1":1,"SI2":2,"SI1":3,"VS2":4,"VS1":5,"VVS2":6,"VVS1":7,"IF":8}
```

```
16) df['color'].unique()
```

```
17) array(['F', 'J', 'G', 'E', 'D', 'H', 'I'], dtype=object)
```

```
18) color_map = {"D":1,"E":2,"F":3,"G":4,"H":5,"I":6,"J":7}
```

```
19) df['cut']=df['cut'].map(cut_map)
    df['clarity']=df['clarity'].map(clarity_map)
    df['color']=df['color'].map(color_map)
```

```
20) df.head()
```

```
21)
   carat  cut  color  clarity  depth  table    x     y     z  price
0    1.52  4     G      SI2     62.2   58.0   7.27  7.33  4.55  13619
1    2.03  3     F      SI1     62.0   58.0   8.06  8.12  5.05  13387
2    0.70  5     G      SI1     61.2   57.0   5.69  5.73  3.90   2772
3    0.52  5     G      SI1     61.6   56.0   4.38  4.41  2.71    666
4    1.70  4     G      SI1     62.6   59.0   7.65  7.61  4.77  14453
```

Model Training

```
[6]: df=pd.read_csv('https://raw.githubusercontent.com/krishnaik86/FSDSRegression/main/notebooks/data/gaestone.csv')
df.head()
```

```
[8]:
```

	id	carat	cut	color	clarity	depth	table	x	y	z	price
0	0	1.52	Premium	F	VS2	62.2	58.0	7.27	7.33	4.55	13619
1	1	2.03	Very Good	J	S12	62.0	58.0	8.06	8.12	5.05	13387
2	2	0.70	Ideal	G	VS1	61.2	57.0	5.99	5.73	3.50	2772
3	3	0.32	Ideal	G	VS1	61.6	56.0	4.38	4.41	2.71	666
4	4	1.70	Premium	G	VS2	62.6	59.0	7.65	7.61	4.77	14453

```
[9]: df=df.drop(labels=['id'],axis=1)
```

```
[10]: ## Independent and dependent features
X = df.drop(labels=['price'],axis=1)
Y = df[['price']]
```

```
[11]: Y
```

```
[11]:
```

	price
0	13619
1	13387
2	2772
3	666
4	14453
...	...
193568	1130
193569	2874
193570	3096
193571	681
193572	2258

193573 rows x 1 columns

[12]:

X

[12]:

	carat	cut	color	clarity	depth	table	x	y	z
0	1.52	Premium	F	VS2	62.2	58.0	7.27	7.33	4.55
1	2.03	Very Good	F	S12	62.0	58.0	8.06	8.12	5.05
2	0.70	Ideal	G	VS1	61.2	57.0	5.68	5.73	3.50
3	0.32	Ideal	G	VS1	61.6	56.0	4.38	4.41	2.71
4	1.70	Premium	G	VS2	62.6	59.0	7.65	7.61	4.77
...
193568	0.31	Ideal	D	VS2	61.1	56.0	4.35	4.39	2.67
193569	0.70	Premium	G	VS2	60.3	58.0	5.75	5.77	3.47
193570	0.73	Very Good	F	S11	63.1	57.0	5.72	5.75	3.62
193571	0.34	Very Good	D	S11	62.9	55.0	4.45	4.49	2.81
193572	0.71	Good	E	S12	60.8	64.0	5.73	5.71	3.48

193573 rows x 9 columns

[13]:

```
# Define which columns should be ordinal-encoded and which should be scaled
categorical_cols = X.select_dtypes(include='object').columns
numerical_cols = X.select_dtypes(exclude='object').columns
```

[14]:

```
# Define the custom ranking for each ordinal variable
cut_categories = ['Fair', 'Good', 'Very Good', 'Premium', 'Ideal']
color_categories = ['D', 'E', 'F', 'G', 'H', 'I', 'J']
clarity_categories = ['I1', 'SI2', 'SI1', 'VS2', 'VS1', 'VS2', 'VS1', 'IF']
```

[15]:

```
from sklearn.impute import SimpleImputer # Handling Missing Values
from sklearn.preprocessing import StandardScaler # Handling Feature Scaling
from sklearn.preprocessing import OrdinalEncoder # Ordinal Encoding
## pipelines
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
```

```

6) # Numerical Pipeline
num_pipeline=Pipeline(
    steps=[
        ('imputer',SimpleImputer(strategy='median')),
        ('scaler',StandardScaler())
    ]
)

# Categorical Pipeline
cat_pipeline=Pipeline(
    steps=[
        ('imputer',SimpleImputer(strategy='most_frequent')),
        ('ordinalencoder',OrdinalEncoder(categories=[cut_categories,color_categories,clarity_categories])),
        ('scaler',StandardScaler())
    ]
)

preprocessor=ColumnTransformer([
    ('num_pipeline',num_pipeline,numerical_cols),
    ('cat_pipeline',cat_pipeline,categorical_cols)
])

```

```

7) # train test split

from sklearn.model_selection import train_test_split

X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.30,random_state=30)

```

```

8) X_train=pd.DataFrame(preprocessor.fit_transform(X_train),columns=preprocessor.get_feature_names_out())
X_test=pd.DataFrame(preprocessor.transform(X_test),columns=preprocessor.get_feature_names_out())

```

```

9) X_train.head()

```

```

10)

```

	num_pipeline_carat	num_pipeline_depth	num_pipeline_table	num_pipeline_x	num_pipeline_y	num_pipeline_z	cat_pipeline_cut
0	0.975439	0.849607	-0.121531	-1.042757	-1.080970	-1.123150	0.874076
1	0.235195	1.833637	-0.121531	0.318447	0.275859	0.485354	2.144558
2	0.494617	0.815855	0.399800	0.570855	0.606458	0.673737	0.132136
3	1.018676	0.260701	0.921131	-1.214034	-1.244270	-1.195605	0.132136
4	0.953821	0.664555	0.642862	-1.069801	-1.044681	-1.094168	0.874076

```
1: X_test.head()
```

```
1: 
```

	num_pipeline_carat	num_pipeline_depth	num_pipeline_table	num_pipeline_x	num_pipeline_y	num_pipeline_z	cat_pipeline_cut
0	0.564688	0.942132	0.642862	0.429765	0.464061	0.500096	0.132136
1	0.175556	1.000906	0.121531	0.042137	0.028595	0.036132	1.138347
2	1.061913	0.260701	0.121531	1.304180	1.298703	1.268060	0.874076
3	0.970223	0.201927	1.963794	1.048629	0.996563	0.978046	0.132136
4	0.982202	1.312235	0.399800	1.006699	0.990248	1.065186	0.132136

```
1: ## Model Training
```

```
from sklearn.linear_model import LinearRegression,Lasso,Ridge,ElasticNet
from sklearn.metrics import r2_score,mean_absolute_error,mean_squared_error
```

```
1: regression=LinearRegression()
regression.fit(X_train,y_train)
```

```
1: LinearRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
1: regression.coef_
```

```
1: array([[ 6433.66883504, -132.75843566, -78.42922179, -1728.38971463,
        -499.20382619, -63.39317848,  72.44537247, -468.41684642,
         658.76431651]])
```

```
1: regression.intercept_
```

```
1: array([3978.76628955])
```

```
1: import numpy as np
def evaluate_model(true, predicted):
    mae = mean_absolute_error(true, predicted)
    mse = mean_squared_error(true, predicted)
    rmse = np.sqrt(mean_squared_error(true, predicted))
    r2_square = r2_score(true, predicted)
    return mae, rmse, r2_square
```

```

1: ## Train multiple models

models=[
    'LinearRegression':LinearRegression(),
    'Lasso':Lasso(),
    'Ridge':Ridge(),
    'ElasticNet':ElasticNet()
]
trained_model_list=[]
model_list=[]
r2_list=[]

for i in range(len(list(models))):
    model=list(models.values())[i]
    model.fit(X_train,y_train)

    ## Make Predictions
    y_pred=model.predict(X_test)

    mae, rmse, r2_square=evaluate_model(y_test,y_pred)

    print(list(models.keys())[i])
    model_list.append(list(models.keys())[i])

    print('Model Training Performance')
    print("RMSE:",rmse)
    print("MAE:",mae)
    print("R2 score",r2_square*100)

    r2_list.append(r2_square)

    print('='*35)
    print('\n')

```

```

LinearRegression
Model Training Performance
RMSE: 1813.9847894344882
MAE: 674.825511579685
R2 score 93.68988248567512
*****

```

```

Lasso
Model Training Performance
RMSE: 1813.8784226767813
MAE: 675.871602336216
R2 score 93.68948971841784
*****

```

```

Ridge
Model Training Performance
RMSE: 1813.9859272771631
MAE: 674.8555888798284
R2 score 93.6898673258594
*****

```

```

Elasticnet
Model Training Performance
RMSE: 1533.4163456864848
MAE: 1868.7368759154729
R2 score 85.56494831165182
*****

```

```

1: model_list

```

```

1: ['LinearRegression', 'Lasso', 'Ridge', 'Elasticnet']

```